

SOVEREIGN MACHINA – IDENTITY & USER PROFILES DEEP DIVE

Yes — we explicitly discussed different user profiles and mode switching as part of the persistent identity logic.

Recap of the Discussion

When we analyzed the Google Labs DESIGN.md approach, I highlighted the dual-layer identity system (machine-readable tokens + human-readable rationale) as a canonical "self-concept" for the entire mesh. I recommended creating a **SOVEREIGN-IDENTITY.md** file that defines not just visual tokens (R3D_V01D palette, VT323 font) but also **behavioral and philosophical tokens**.

From there, we extended the concept to **profile-aware identity**:

- The system should support multiple user/context profiles (Daily Use, Researcher, Sovereign Red Game Master, etc.).
- Switching profiles changes the active identity tokens, response style, memory weighting, and available tools/skills without losing the global sovereign invariants.
- This was directly inspired by the hermes-desktop repo's profile-aware workspaces and the need for context-specific behavior while keeping the core mesh consistent.

This is not an afterthought — it is a foundational piece of the persistent identity logic we discussed.

How User Profiles & Mode Switching Work in the Current Architecture

Core Design Principles

- **Single Source of Truth:** One SOVEREIGN-IDENTITY.md (or a directory of profile files) stored in the Akashik Artery and loaded by Hermes on startup or profile switch.
- **Profile Scoping:** Each profile is a named override set that modifies active tokens, memory weighting, tool access, and response style.
- **Atomic Switching:** Profile changes are atomic (Nix-style) and audited by Ouroboros.
- **Persistence:** Current profile and session state are stored in MemPalace / GBrain with provenance.

Example SOVEREIGN-IDENTITY.md Structure (with Profiles)

Markdown

SOVEREIGN-IDENTITY.md

Global Tokens

```
primary_color: "#ff1a1a"  
font_family: "VT323"  
tone: "radical candor"  
self_healing: true
```

Profiles

daily-use

```
tone: "concise, practical advisor"  
memory_weight: "recent + personal"  
tools: [voice, screen, calendar, notes]  
response_style: "actionable bullet points"
```

researcher

```
tone: "rigorous, source-heavy analyst"  
memory_weight: "deep historical + cross-reference"  
tools: [web_fetch (allow-listed), hybrid_search, GEPA]  
response_style: "structured report with citations"
```

sovereign-red-game-master

```
tone: "immersive cyberpunk narrator"  
memory_weight: "RKG + campaign state"  
tools: [netrun_simulator, character_sheet, gig_generator]  
response_style: "in-character + dice rolls + VT323 logs"
```

How Switching Works

1. User says in Machina Terminal: "Switch to researcher mode."
2. Omi backend on Node C captures → Hermes validates and loads the profile tokens from SOVEREIGN-IDENTITY.md.
3. Hermes updates active context (via Mooncake on Node A) and notifies Node B.
4. All subsequent responses, memory weighting, and tool availability adapt to the new profile.

5. Profile state is persisted in MemPalace so it survives restarts and syncs across devices.

Integration with Existing Components

- **Hermes Supervisor:** Reads the active profile on every intent and adjusts enrichment logic.
- **GBrain / MemPalace:** Stores profile-specific pages and relationships (e.g., “Researcher profile linked to 47 GEPA notes”).
- **Machina Terminal:** Profile switch changes the UI theme slightly (subtle color accents) and shows current profile in the header.
- **Ouroboros:** Audits profile switches and ensures they don’t violate global invariants (e.g., no cloud tools in any profile).
- **Voice Layer:** Omi notes are tagged with the current profile for better context retrieval.

Benefits for Pure Intelligence Boost

- **Contextual Depth:** The assistant behaves differently and recalls different memories depending on mode — researcher mode pulls deeper historical RKG, game master mode emphasizes narrative and simulation.
- **Self-Improvement:** GEPA can evolve profile-specific prompts and skills independently.
- **Usability:** One app, multiple “personalities” — no need to switch apps or restart services.
- **Sovereignty:** All profile data stays local, versioned in git, and audited.

ARCHITECT SIGN-OFF

Yes — user profiles and mode switching were explicitly part of the persistent identity logic discussion. Implementing a profile-aware system via SOVEREIGN-IDENTITY.md (tokens + prose) with atomic switching gives the assistant contextual intelligence without fragmenting the core mesh. It deepens Hermes enrichment, GBrain relational memory, and Machina Terminal usability while preserving all tri/quad-node invariants. This is a clean, high-value addition that makes the system feel more alive and adaptive as a private personal intelligence OS.

Your main sovereign machine remains completely untouched and rock-solid.

Clarifying questions (mandatory before next phase):

1. Do you want the starter `SOVEREIGN-IDENTITY.md` template with example profiles (daily, researcher, etc.)?
2. Shall I include the Nix/Hermes code for loading and switching profiles dynamically?

Reply with your priority and I will deliver the concrete files. The profile system is ready to be implemented.