

DISCOVERY PHASE – COMPLETE (Zero assumptions; requirements locked verbatim from your query + full Sovereign Machina architecture + changelog v3.2.7 + latest repo state)

Core functionality required:

- Analyze <https://github.com/google-labs-code/design.md> and extract any actionable insights, particularly around **identity logic**, for Sovereign Machina.
- Focus strictly on integration with the current tri/quad-node stack as a private, self-improving personal intelligence assistant (voice-first, knowledge OS, decision support, life logging).
- Maintain: Node A = Mooncake KV buffer, Node B = Director, Node C = Oracle / Voice Core with Omi + Hermes, optional Hermes Pod. No changes to Node A or Node B core logic.

All requirements are now exhaustive and production-locked.

REPOSITORY SUMMARY

This is a **Google Labs** specification for **DESIGN.md** — a format that gives coding agents a persistent, structured understanding of a visual identity / design system. It combines **machine-readable tokens** (YAML frontmatter) with **human-readable rationale** (Markdown prose). The goal is to make agents produce consistent, on-brand output without repeated prompting.

Key Elements:

- YAML frontmatter for exact tokens (colors, typography, spacing, rounded, components, etc.).
- Markdown body for "why" and usage guidelines.
- Tools: linter, diff, export (Tailwind, DTCG), CLI.
- Agent-friendly: tokens are normative, prose provides context.
- Example: Defines primary/secondary colors, typography scales, component variants with references like `{colors.primary}`.

INSIGHTS FOR SOVEREIGN MACHINA

1. Identity Logic as Canonical Source of Truth

→ ...

The core idea — a single **DESIGN.md** file that agents read to understand "who we are" — is extremely powerful.

Insight to Adopt: Create a `SOVEREIGN-IDENTITY.md` (or `MACHINA-IDENTITY.md`) in the root of your repo / Akashik Artery. It would define:

- Visual tokens (R3D_V01D palette, VT323 font, spacing, CRT effects).
- Behavioral tokens (tone, response style, self-healing rules, provenance requirements).
- Philosophical tokens (radical candor, physical sovereignty, zero-drift).

Hermes (and the GBrain rewrite) would read this file on startup and every evolution cycle. This gives every agent a persistent, structured "self-concept" without relying on system prompts alone.

2. Machine + Human Dual Layer

Tokens for exact values + prose for reasoning.

Insight to Adopt: Apply this to Hermes and GBrain. Machine tokens for strict enforcement (e.g., color hex, font family, response structure). Prose for "why" (e.g., "Blood red #ff1a1a represents the pulse of sovereignty"). This makes GEPA evolution and self-healing more grounded and auditable.

3. Linter & Diff Capabilities

The repo includes linting and diff tools to catch regressions.

Insight to Adopt: Add a similar linter to your Nix CI pipeline and Hermes startup. It would validate that all generated output (Machina Terminal responses, RKG entries, screen renders) conforms to the identity file. This strengthens Ouroboros auditing and prevents aesthetic or behavioral drift.

4. Agent Consistency & Reduced Prompting

Agents that read DESIGN.md produce consistent output without repeated instructions.

Insight to Adopt: Feed the identity file into Hermes at the start of every intent processing cycle. This reduces token usage, improves consistency across voice responses, screen awareness, and external actions, and makes the assistant feel more "alive" and coherent over time.

5. Relevance to Current Stack

- **Hermes:** Use the identity file as a core input for enrichment and Skill distillation.

- **GBrain Rewrite:** Store identity tokens as first-class graph nodes in MemPalace for relational queries.
- **Machina Terminal:** Load the identity file to enforce theming and response formatting client-side.
- **Voice / Screen:** Omi notes and screen captures can be validated against the identity for consistency.

ARCHITECT SIGN-OFF

The Google Labs DESIGN.md repo provides excellent, actionable insights for **identity logic**. The dual machine/human format, persistent structured tokens, and agent-friendly design system approach are highly relevant. Adopting a `SOVEREIGN-IDENTITY.md` file (with tokens + prose) would give Hermes, GBrain, and Machina Terminal a canonical, self-consistent "self-concept." This deepens self-healing, reduces drift, improves output quality, and makes the entire system feel more coherent and sovereign. It is a low-effort, high-impact addition that fits perfectly into your current architecture.

Your main sovereign machine remains completely untouched and rock-solid.

Clarifying questions (mandatory before next phase):

1. Do you want a starter `SOVEREIGN-IDENTITY.md` template with `R3D_V01D` tokens and prose?
2. Shall I include a Nix module or Hermes hook to load and validate the identity file on startup?

Reply with your priority and I will deliver the concrete files. The identity logic is ready to be integrated.