

✓ EVALUATION: GLM Family (Z.AI / Zhipu AI) in Sovereign Machina Topology

Short Answer: Yes — worth evaluating seriously, especially as a **strong coding / agentic specialist** on Node B or Node D. GLM models are among the top open-weight performers in 2026 for real-world coding, tool-use, and agentic workflows.

Current GLM Landscape (April 2026)

- **GLM-4.7 / GLM-4.7-Flash** (30B-class MoE, ~3B active) — Frequently called one of the best local coding/agentic models. Strong on SWE-bench, τ^2 -Bench, terminal tasks, and frontend generation.
- **GLM-4.6 / GLM-4.5-Air** — Excellent balance of reasoning + coding.
- Strong vision variants (GLM-4.6V / GLM-4.5V) with native multimodal support.
- Good GGUF support via Unsloth, bartowski, etc.
- Apache/MIT-style licensing in many releases.

Strengths vs Current Stack:

- Often beats or matches Qwen Coder on practical coding/agentic tasks.
- Excellent tool-calling and multi-turn agent behavior.
- Strong in "thinking before acting" mode.
- Competitive with Gemma-4 26B in many agentic scenarios.

Weaknesses:

- Slightly less consistent tokenizer/tool-calling compared to Gemma-4 family (may require extra prompt engineering).
- Community support / GGUF ecosystem is good but smaller than Gemma/Qwen.
- Some versions have chat template quirks in llama.cpp.

Recommendation for Your Topology

GLM is worth adding as a specialist, but not replacing the core Gemma-4 26B.

Proposed Integration (Optional but High-Value):

- **Node B (Director, 16 GB):** Keep Gemma-4 26B Q3 (as you chose) for orchestration. Add GLM-4.7-Flash Q5/Q6 (~12–15 GB) as **secondary coding/agentive specialist** if headroom allows.
- **Node D (K15, 48 GB):** Primary Gemma-4 26B Q6 + Qwen Coder 14B. Optionally add a GLM-4.7-Flash variant for diversity in agentive tasks.

This gives you **three strong coding perspectives** (Gemma generalist + Qwen fast coder + GLM agentive specialist) while keeping the topology clean.

Would you like me to:

1. Add GLM-4.7-Flash to the model farm (with exact GGUF links and Nix config)?
2. Compare specific GLM variant vs your current plan in detail?
3. Keep current plan (Gemma + Qwen only) and finalize Nix configs?

Reply with your preference and we'll lock the final topology.