

## ✅ LOCKED RECOMMENDATION FOR NODE B

Use `mmproj-BF16.gguf` (1.19 GB)

### Why BF16 is the best choice for Node B

- **Gemma-4 models** (including the 26B A4B) were **trained in BF16**. Using the matching `mmproj-BF16.gguf` preserves the original precision and gives the best quality/speed balance.
- Your **9060 XT (RDNA3)** has excellent BF16 support — no casting penalties.
- **F16.gguf** is acceptable but can introduce slight quality loss on Gemma models (some users report worse vision coherence).
- **F32.gguf** (2.29 GB) is overkill — higher quality but wastes VRAM and offers almost no real-world benefit for vision tasks on this setup.

**Rule of thumb (2026 community consensus):**

**Always prefer BF16 mmproj** when available for Gemma-4 vision. It's the sweet spot.

### Final Node B Vision Setup

- Main model: **Gemma-4 26B A4B Q3\_K\_M** (your choice)
- Multimodal projector: `mmproj-BF16.gguf`

This keeps VRAM usage low while giving solid vision capabilities (screenshots, diagrams, Obsidian canvas analysis, etc.) to the Director layer.

**Would you like me to:**

1. Give you the **exact llama.cpp / vLLM launch command** for Node B with this mmproj?
2. Update the full **Nix model-farm service** for Node B including vision support?
3. Both + integration notes for Pretext HUD / Obsidian?

Just say the word and we'll finalize Node B completely.