

LOCKED: FINAL CLEAN MODEL TOPOLOGY v1.1

User Decision: Node B = Gemma-4 26B Q3_K_M (clean)

I have updated the entire cluster plan based on your choice.

FINALIZED MODEL TOPOLOGY (Clean Bases Only)

Node	Hardware	Role	Model(s)	Quant	Approx. Usage	Headroom	Purpose
Node A	Nitro 5 (1050 Ti 4 GB)	Mooncake KV Buffer	(No inference)	—	—	Excellent	Shared long-term context & MemPalace
Node B	9060 XT (16 GB)	Director / Orchestrator	Gemma-4 26B A4B (clean)	Q3_K_M	~11–13 GB	3–5 GB	Fast Director, swarm coordination, profile switching, medium synthesis, Pretext HUD logic
Node C	RTX 2060 (6 GB)	Fast / Light Tier	Gemma-4-E4B (clean) Gemma-4-E4B (clean) Falcon Perception	Q3_K_M Q4_K_M Q4/Q5	~3.8 GB ~4.8 GB ~1.2 GB	~0.5–1 GB	Ultra-fast routing, simple intent, light synthesis, basic vision
Node D	GMKtec K15 (48 GB)	Hermes Command Center	Gemma-4 26B A4B (clean) Qwen2.5-Coder-14B (clean)	Q6_K Q6_K	~23.5 GB ~10.5 GB	12–14 GB	Heavy agentic reasoning, GEPA v2 evolution, complex planning, fast coding specialist

Summary of Changes & Rationale

- **Node B:** You opted for **Q3_K_M** on the 26B model.
 - This gives excellent speed and **3–5 GB headroom** on 16 GB VRAM.
 - Still significantly stronger than the old E4B Q8 while staying very responsive as Director.
 - Perfect balance for orchestration and swarm coordination.
- **Node C:** Kept as pure fast tier (Q3 + Q4 + Falcon). No changes.
- **Node D:** Remains heavy tier with 26B Q6 + Qwen Coder 14B Q6 (strong headroom for Hermes + GEPA).
- All models are now **clean base** only.
- GEPA v2 + Hyper-Tune weighting will create specialized Director, Researcher, Dev-Mode, etc. genomes on top of these clean bases.

Updated Routing Logic (hermes-inference-router)

```
Rust

pub async fn select_backend(intent: &Intent, profile: &Profile) -> Backend
{
    match intent.task_type {
        TaskType::UltraFast           => Backend::NodeC("gemma-q3"),
        TaskType::MediumSynthesis      => Backend::NodeC("gemma-q4") |
Backend::NodeB("gemma-26b-q3"),
        TaskType::FastCoding           => Backend::NodeD("qwen-coder-14b"),
        TaskType::HeavyReasoning       => Backend::NodeD("gemma-26b-q6"),
        TaskType::ComplexAgentic       => Backend::NodeD("gemma-26b-q6"),
        TaskType::Orchestration        => Backend::NodeB("gemma-26b-q3"),
        // Director
        _                               => Backend::NodeB("gemma-26b-q3"),
    }
}
```

This is now a **clean, balanced, and stable** topology with good headroom on every node.

Next Steps – Ready to Deploy

Reply with the number(s) you want:

1. Full Nix `model-farm` configs for **all nodes** (download + launch commands)
2. Updated `hermes-inference-router.rs` + profile mappings
3. GEPA v2 config tuned for clean bases
4. Migration script to purge old Obliterated models
5. All of the above (complete cleanup package)

We can now finalize the cluster and move full development inside Sovereign Machina with confidence.

What would you like next?