✅ **LOCKED DECISION: Cut Obliterated Models – Switch to Clean Bases + Hyper-Tune Weighting**

Yes. You are making the correct call.

Why Cut the Obliterated Models?

Current Situation on Node C (and partially Node D)

- You are running "Gemma-4-E4B-it-OBLITERATED" variants (heavily modified, uncensored, agentic-tuned).
- These are great for raw creativity and "jailbreak-style" behavior, but they introduce **inconsistencies** in tool-calling, structured output, reasoning chains, and GEPA evolution.
- They make provenance harder and can cause drift in Hermes self-evolution.

Recommendation: Purge all Obliterated variants and move to **clean base models** + your own **Hyper-Tune weighting** (via GEPA v2 + DSPy).

This is the professional, production-grade path for a self-improving sovereign system.

New Locked Model Strategy

Node C (6 GB – Fast Tier) – Clean Only

- Gemma-4-E4B **Q3_K_M** (clean) → Ultra-fast routing
- Gemma-4-E4B **Q4_K_M** (clean) → Light-medium synthesis
- Falcon Perception (clean/small)

Node D (K15 48 GB – Heavy Tier) – Clean Only

- **Gemma-4 26B A4B Q6_K** (clean base) → Primary reasoning + GEPA evolution
- **Qwen2.5-Coder-14B Q6_K** (clean base) → Dedicated fast coding

Node B remains Gemma-4 Q8 (clean).

Hyper-Tune Weighting Strategy (Your Custom Evolution Layer)

Instead of relying on pre-obliterated models, we let **Hermes GEPA v2** do the heavy lifting:

- Start with clean base models (best coherence, tool-calling, and long-context stability).
- Use GEPA reflective evolution + DSPy MIPRO to create **profile-specific hyper-tuned weights** on top of the clean bases.
- Store evolved genomes in MemPalace (per-profile).
- Periodically merge the best performing genomes into lightweight LoRA adapters (if needed) or just keep them as prompt + few-shot genomes.

This gives you **better control** than any pre-obliterated model, because the tuning is **your data + your trajectories**.

Benefits:

- More consistent agentic behavior
- Better GEPA self-evolution (clean base = cleaner learning signal)
- Easier provenance and auditing
- Future-proof (you control the evolution)

Immediate Action Plan

1. **Purge Obliterated models** from all model-farm configs.
2. **Download clean bases** (if not already present).
3. **Update hermes-inference-router** to always prefer clean models.
4. **Boost GEPA v2** to run more aggressive evolution on clean bases (increase population to 16 on Node D).

Would you like me to:

1. Give you the **full Nix model-farm refactor** for Node C + Node D (clean models only)?
2. Update the **hermes-inference-router + GEPA config** for clean + hyper-tune preference?
3. Both + a migration script to purge old obliterated files?

Reply with 1, 2, or 3 and I'll deliver the exact diffs right now.

This is the right architectural move for long-term stability and self-improvement. Clean bases + your own hyper-tuning will outperform pre-obliterated models in the long run.

Let’s clean it up and make it yours. Your choice?