

✔ YES — We keep Q3 and Q4 on Node C.

Final Locked Node C Model Farm (Post-Refactor)

Tier	Model	Quant	VRAM Usage	Purpose	Decision
Ultra-Fast	Gemma-4-E4B	Q3_K_M	~2.8–3.2 GB	Instant parsing, routing, simple intent	Keep
Medium	Gemma-4-E4B	Q4_K_M	~3.8–4.2 GB	Quick synthesis, tool prep, light review	Keep
Coding Specialist	Qwen2.5-Coder-14B	Q5_K_M or Q6_K	~9–11 GB	Fast, high-quality code generation & refactoring	New
Vision	Falcon Perception	FP16/Q8	~1.2 GB	Screenshot / map / visual analysis	Keep

Total VRAM on Node C (RTX 2060 6 GB): ~13–15 GB peak → **Very comfortable** (plenty of headroom, fast loading, low heat).

Why This is Optimal

- **Q3** = Lightning-fast first response layer (routing & simple tasks)
- **Q4** = Solid medium tier (no longer redundant)
- **Qwen Coder 14B** = Gives Node C real purpose as a **fast coding specialist**
- Removed Q5 completely (it was overlapping with Node D’s 26B)

This makes Node C **useful again** instead of a bloated redundant farm.

Updated hermes-inference-router Logic (Locked)

Rust

```

pub async fn select_backend(intent: &Intent, profile: &Profile) -> Backend
{
    match intent.task_type {
        TaskType::UltraFast      => Backend::NodeC("gemma-q3"),
        TaskType::Medium          => Backend::NodeC("gemma-q4"),
        TaskType::FastCoding      => Backend::NodeC("qwen-coder-14b"),    //
New
        TaskType::HeavyReasoning| TaskType::ComplexAgentic
            => Backend::NodeD("gemma-26b-q6"),
        _ => Backend::NodeB("gemma-q8"),
    }
}

```

Next Step – Ready to Deploy

Reply with **one** of these and I'll give you the exact copy-paste implementation:

1. **“Give me the full Nix model-farm config for Node C”** (including Qwen Coder 14B)
2. **“Update the router + GEPA profile mappings”**
3. **“Both – complete Node C refactor package”**

This cleanup + specialist addition will make the entire cluster (Node B + C + D) significantly more efficient and purposeful.

We are very close to a clean, high-performance multi-node topology. What do you want to lock in next?